

```
%matplotlib inline ### import libraries import numpy as np import pandas as pd import matplotlib.pyplot as plt from matplotlib import style
import seaborn as sns
```

```
In [2]:
```

```
data = pd.read_csv('health care diabetes.csv')
```

```
In [3]:
```

```
data.head()
```

```
Out[3]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 | 1 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 | 0 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 | 1 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 | 0 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 | 1 |

```
In [4]:
```

```
data.isnull().any()
```

```
Out[4]:
```

```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction False
Age               False
Outcome           False
dtype: bool
```

```
In [5]:
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose          768 non-null int64
BloodPressure    768 non-null int64
SkinThickness    768 non-null int64
Insulin          768 non-null int64
BMI              768 non-null float64
DiabetesPedigreeFunction 768 non-null float64
Age              768 non-null int64
Outcome          768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [41]:
```

```
Positive = data[data['Outcome']==1]
Positive.head(5)
```

```
Out[41]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|--|-------------|---------|---------------|---------------|---------|-----|--------------------------|-----|---------|
|--|-------------|---------|---------------|---------------|---------|-----|--------------------------|-----|---------|

| 0 | Pregnancies ⁶ | Glucose ¹⁴⁸ | BloodPressure ⁷² | SkinThickness ³⁵ | Insulin ⁰ | BMI ^{23.6} | DiabetesPedigreeFunction ^{0.627} | Age ⁵⁰ | Outcome ¹ |
|---|--------------------------|------------------------|-----------------------------|-----------------------------|----------------------|---------------------|---|-------------------|----------------------|
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 | 1 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 | 1 |
| 6 | 3 | 78 | 50 | 32 | 88 | 31.0 | 0.248 | 26 | 1 |
| 8 | 2 | 197 | 70 | 45 | 543 | 30.5 | 0.158 | 53 | 1 |

In [43]:

```
data['Glucose'].value_counts().head(7)
```

Out[43]:


```
100    17
99     17
129    14
125    14
111    14
106    14
95     13
Name: Glucose, dtype: int64
```

In [35]:

```
plt.hist(data['Glucose'])
```

Out[35]:

```
(array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
 array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
        179.1, 199. ]),
 <a list of 10 Patch objects>)
```



In [33]:

```
data['BloodPressure'].value_counts().head(7)
```

Out[33]:

```
70     57
74     52
68     45
78     45
72     44
64     43
80     40
Name: BloodPressure, dtype: int64
```


In [36]:

```
plt.hist(data['BloodPressure'])
```

Out[36]:

```
(array([ 25,  50,  75, 100, 125, 150, 175, 200, 225, 250, 275, 300, 325, 350, 375, 400, 425, 450, 475, 500, 525, 550, 575, 600, 625, 650, 675, 700, 725, 750, 775, 800, 825, 850, 875, 900, 925, 950, 975, 1000]),
```

```
(array([ 35.,  1.,  2., 13., 107., 261., 243.,  87., 14.,  5.]),
 array([ 0., 12.2, 24.4, 36.6, 48.8, 61., 73.2, 85.4, 97.6,
        109.8, 122. ]),
 <a list of 10 Patch objects>)
```



In [32]:

```
data['SkinThickness'].value_counts().head(7)
```

Out[32]:


```
0      227
32     31
30     27
27     23
23     22
33     20
18     20
Name: SkinThickness, dtype: int64
```

In [37]:

```
plt.hist(data['SkinThickness'])
```

Out[37]:

```
(array([231., 107., 165., 175.,  78.,  9.,  2.,  0.,  0.,  1.]),
 array([ 0.,  9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
 <a list of 10 Patch objects>)
```



In [30]:

```
data['Insulin'].value_counts().head(7)
```

Out[30]:

```
0      374
105     11
140      9
130      9
120      8
```


```
100      7
94       7
Name: Insulin, dtype: int64
```

In [38]:

```
plt.hist(data['Insulin'])
```

Out[38]:

```
(array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
 array([ 0., 84.6, 169.2, 253.8, 338.4, 423., 507.6, 592.2, 676.8,
        761.4, 846. ]),
 <a list of 10 Patch objects>)
```



In [29]:

```
data['BMI'].value_counts().head(7)
```

Out[29]:


```
32.0    13
31.6    12
31.2    12
0.0     11
33.3    10
32.4    10
32.8     9
Name: BMI, dtype: int64
```

In [39]:

```
plt.hist(data['BMI'])
```

Out[39]:

```
(array([ 11.,  0., 15., 156., 268., 224., 78., 12., 3., 1.]),
 array([ 0., 6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
        60.39, 67.1 ]),
 <a list of 10 Patch objects>)
```



In [9]:

```
data.describe().transpose()
```

Out[9]:

| | count | mean | std | min | 25% | 50% | 75% | max |
|---------------------------------|-------|------------|------------|--------|----------|----------|-----------|--------|
| Pregnancies | 768.0 | 3.845052 | 3.369578 | 0.000 | 1.00000 | 3.0000 | 6.00000 | 17.00 |
| Glucose | 768.0 | 120.894531 | 31.972618 | 0.000 | 99.00000 | 117.0000 | 140.25000 | 199.00 |
| BloodPressure | 768.0 | 69.105469 | 19.355807 | 0.000 | 62.00000 | 72.0000 | 80.00000 | 122.00 |
| SkinThickness | 768.0 | 20.536458 | 15.952218 | 0.000 | 0.00000 | 23.0000 | 32.00000 | 99.00 |
| Insulin | 768.0 | 79.799479 | 115.244002 | 0.000 | 0.00000 | 30.5000 | 127.25000 | 846.00 |
| BMI | 768.0 | 31.992578 | 7.884160 | 0.000 | 27.30000 | 32.0000 | 36.60000 | 67.10 |
| DiabetesPedigreeFunction | 768.0 | 0.471876 | 0.331329 | 0.078 | 0.24375 | 0.3725 | 0.62625 | 2.42 |
| Age | 768.0 | 33.240885 | 11.760232 | 21.000 | 24.00000 | 29.0000 | 41.00000 | 81.00 |
| Outcome | 768.0 | 0.348958 | 0.476951 | 0.000 | 0.00000 | 0.0000 | 1.00000 | 1.00 |

In []:


Week 2

In [49]:

```
plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

Out[49]:

```
(array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
        14.,  7.,  3.,  1.,  1.,  0.,  1.]),
array([ 0.,  3.355,  6.71, 10.065, 13.42, 16.775, 20.13, 23.485,
        26.84, 30.195, 33.55, 36.905, 40.26, 43.615, 46.97, 50.325,
        53.68, 57.035, 60.39, 63.745, 67.1 ]),
<a list of 1 Patch objects>)
```



In [55]:

```
Positive['BMI'].value_counts().head(7)
```

Out[55]:

```
32.9    8
31.6    7
33.3    6
30.5    5
32.0    5
31.2    5
```


```
31.2      3
32.4      4
Name: BMI, dtype: int64
```

In [61]:

```
plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

Out[61]:

```
(array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
array([  0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199.  ]),
<a list of 1 Patch objects>)
```



In [56]:

```
Positive['Glucose'].value_counts().head(7)
```

Out[56]:


```
125      7
158      6
128      6
115      6
129      6
146      5
162      5
Name: Glucose, dtype: int64
```


In [62]:

```
plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

Out[62]:

```
(array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
        36., 41., 31.,  7.,  4.,  4.,  3.]),
array([  0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
        51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
        102.6, 108.3, 114.  ]),
<a list of 1 Patch objects>)
```





In [57]:

```
Positive['BloodPressure'].value_counts().head(7)
```

Out[57]:


```
70    23
76    18
78    17
74    17
72    16
0      16
82     13
Name: BloodPressure, dtype: int64
```

In [63]:

```
plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

Out[63]:

```
(array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
        0.,  0.,  0.,  0.,  0.,  0.,  1.]),
 array([ 0.,  4.95,  9.9, 14.85, 19.8, 24.75, 29.7, 34.65, 39.6,
        44.55, 49.5, 54.45, 59.4, 64.35, 69.3, 74.25, 79.2, 84.15,
        89.1, 94.05, 99. ]),
 <a list of 1 Patch objects>)
```



In [60]:

```
Positive['SkinThickness'].value_counts().head(7)
```

Out[60]:

```
0      88
32     14
33      9
30      9
39      8
35      8
36      8
Name: SkinThickness, dtype: int64
```

In [64]:

```
plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```


Out[64]:

```
(array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
        5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
 array([ 0., 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
```

```

380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
761.4, 803.7, 846. ]),
<a list of 1 Patch objects>)

```



In [59]:

```
Positive['Insulin'].value_counts().head(7)
```

Out[59]:

```

0      138
130      6
180      4
156      3
175      3
194      2
125      2
Name: Insulin, dtype: int64

```

In [65]:

```
#Scatter plot
```

In [68]:

```

BloodPressure = Positive['BloodPressure']
Glucose = Positive['Glucose']
SkinThickness = Positive['SkinThickness']
Insulin = Positive['Insulin']
BMI = Positive['BMI']


```

In [85]:

```


plt.scatter(BloodPressure, Glucose, color=['b'])
plt.xlabel('BloodPressure')
plt.ylabel('Glucose')
plt.title('BloodPressure & Glucose')
plt.show()

```




In [101]:

```
g = sns.scatterplot(x= "Glucose" ,y= "BloodPressure",  
                    hue="Outcome",  
                    data=data);
```




In [100]:

```
B = sns.scatterplot(x= "BMI" ,y= "Insulin",  
                    hue="Outcome",  
                    data=data);
```



In [107]:

```
S = sns.scatterplot(x= "SkinThickness" ,y= "Insulin",  
                    hue="Outcome",  
                    data=data);
```



In [104]:

```
### correlation matrix
```

```
data.corr()
```

Out[104]:


| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|--------------------------|-------------|----------|---------------|---------------|----------|----------|--------------------------|----------|
| Pregnancies | 1.000000 | 0.129459 | 0.141282 | -0.081672 | 0.073535 | 0.017683 | -0.033523 | 0.544341 |
| Glucose | 0.129459 | 1.000000 | 0.152590 | 0.057328 | 0.331357 | 0.221071 | 0.137337 | 0.263514 |
| BloodPressure | 0.141282 | 0.152590 | 1.000000 | 0.207371 | 0.088933 | 0.281805 | 0.041265 | 0.239528 |
| SkinThickness | -0.081672 | 0.057328 | 0.207371 | 1.000000 | 0.436783 | 0.392573 | 0.183928 | 0.113970 |
| Insulin | -0.073535 | 0.331357 | 0.088933 | 0.436783 | 1.000000 | 0.197859 | 0.185071 | 0.042163 |
| BMI | 0.017683 | 0.221071 | 0.281805 | 0.392573 | 0.197859 | 1.000000 | 0.140647 | 0.036242 |
| DiabetesPedigreeFunction | -0.033523 | 0.137337 | 0.041265 | 0.183928 | 0.185071 | 0.140647 | 1.000000 | 0.033561 |
| Age | 0.544341 | 0.263514 | 0.239528 | -0.113970 | 0.042163 | 0.036242 | 0.033561 | 1.000000 |
| Outcome | 0.221898 | 0.466581 | 0.065068 | 0.074752 | 0.130548 | 0.292695 | 0.173844 | 0.238344 |

In [105]:

```
### create correlation heat map
sns.heatmap(data.corr())
```

Out[105]:

<matplotlib.axes._subplots.AxesSubplot at 0x2278a586278>





In [106]:

```
plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(), annot=True, cmap='viridis') ### gives correlation value
```

Out[106]:

<matplotlib.axes._subplots.AxesSubplot at 0x2278a71d710>






In [116]:

```
plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True) ### gives correlation value
```

Out[116]:

<matplotlib.axes._subplots.AxesSubplot at 0x2278bde9f28>



In [113]:

```
# Logistic Regression and model building
```

```
In [117]:
```

```
data.head(5)
```

```
Out[117]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 | 1 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 | 0 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 | 1 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 | 0 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 | 1 |

```
In [130]:
```

```
features = data.iloc[:, [0,1,2,3,4,5,6,7]].values  
label = data.iloc[:,8].values
```

```
In [136]:
```

```
#Train test split  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(features,  
                                                  label,  
                                                  test_size=0.2,  
                                                  random_state =10)
```

```
In [137]:
```

```
#Create model  
from sklearn.linear_model import LogisticRegression  
model = LogisticRegression()  
model.fit(X_train,y_train)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: De  
fault solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.  
FutureWarning)
```

```
Out[137]:
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                   intercept_scaling=1, max_iter=100, multi_class='warn',  
                   n_jobs=None, penalty='l2', random_state=None, solver='warn',  
                   tol=0.0001, verbose=0, warm_start=False)
```

```
In [138]:
```

```
print(model.score(X_train,y_train))  
print(model.score(X_test,y_test))
```

```
0.7833876221498371  
0.7337662337662337
```

```
In [139]:
```

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(label,model.predict(features))  
cm
```

```
Out[139]:
```

```
array([[452,  48],  
       [106, 140]])
```

```
[126, 142]], dtype=int64)
```

In [140]:

```
from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.90 | 0.84 | 500 |
| 1 | 0.75 | 0.53 | 0.62 | 268 |
| micro avg | 0.77 | 0.77 | 0.77 | 768 |
| macro avg | 0.76 | 0.72 | 0.73 | 768 |
| weighted avg | 0.77 | 0.77 | 0.76 | 768 |

In [141]:


```
#Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.834

Out[141]:

[<matplotlib.lines.Line2D at 0x2278c4fba90>]



In [152]:

```
#Applying Decision Tree Classifier
from sklearn.tree import DecisionTreeClassifier
model3 = DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

Out[152]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
```

```
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
```

In [163]:

```
model3.score(X_train,y_train)
```

Out[163]:

```
0.990228013029316
```

In [164]:

```
model3.score(X_test,y_test)
```

Out[164]:

```
0.7532467532467533
```

In [162]:

```
#Applying Random Forest
from sklearn.ensemble import RandomForestClassifier
model4 = RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

Out[162]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=11, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

In [165]:

```
model4.score(X_train,y_train)
```

Out[165]:

```
0.990228013029316
```

In [166]:

```
model4.score(X_test,y_test)
```

Out[166]:

```
0.7532467532467533
```

In [169]:

```
#Support Vector Classifier

from sklearn.svm import SVC
model5 = SVC(kernel='rbf',
              gamma='auto')
model5.fit(X_train,y_train)
```

Out[169]:

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

In [170]:

In [170]:

```
model5model.score(X_test,y_test).score(X_train,y_train)
```

Out[170]:

1.0

In [171]:

```
model5.score(X_test,y_test)
```

Out[171]:

0.6168831168831169

In [142]:

```
#Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
model2.fit(X_train,y_train)
```

Out[142]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=7, p=2,
                     weights='uniform')
```

In [143]:

```
#Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score


# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```


AUC: 0.836

```
True Positive Rate - [0.      0.06716418 0.23880597 0.44776119 0.60074627 0.75373134
 0.88059701 0.98507463 1.      ], False Positive Rate - [0.      0.      0.02  0.056 0.12  0.248 0.42
 8 0.668 1.      ] Thresholds - [2.      1.      0.85714286 0.71428571 0.57142857 0.42857143
 0.28571429 0.14285714 0.      ]
```

Out[143]:

Text(0, 0.5, 'True Positive Rate')





In [144]:


```
#Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.620 auc=0.728 ap=0.728

Out[144]:

[<matplotlib.lines.Line2D at 0x2278d0052e8>]



In [145]:

```
#Precision Recall Curve for KNN

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
```



```


# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.658 auc=0.752 ap=0.709

Out[145]:

[<matplotlib.lines.Line2D at 0x2278d025908>]



In [167]:

```

#Precision Recall Curve for Decission Tree Classifier

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')


```

f1=0.916 auc=0.966 ap=0.958

Out[167]:

[<matplotlib.lines.Line2D at 0x2278bde2a20>]





In [168]:


```
#Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.914 auc=0.968 ap=0.960

Out[168]:

[<matplotlib.lines.Line2D at 0x2278a747cf8>]



In []: